

# L12B: Concurrent Programming

CS1101S: Programming Methodology

Low Kok Lim

November 3, 2021

# Readings

- [SICP JS, Sec 3.4](#): “Concurrency: Time is of the Essence”

# Outline

- Concurrency
- Applications of Serializers
- Concurrent programming for multiple shared resources
- More synchronization mechanisms
- Classical synchronization problems

# Outline

- **Concurrency**
  - Correctness
  - Serialization
- Applications of Serializers
- Concurrent programming for multiple shared resources
- More synchronization mechanisms
- Classical synchronization problems

# What is Concurrent Computing?

- **Concurrent computing**

- A form of computing in which several computations are executed **concurrently** instead of **sequentially** [Wikipedia]
  - **concurrently** — during overlapping time periods
  - **sequentially** — one completing before the next starts

# Why Concurrent Computing?

- **Performance**

- Dividing a complex task into multiple smaller subtasks and executing the subtasks simultaneously on different processors

- **Responsiveness**

- Providing a responsive user interface even when there are other tasks executing independently

- **Abstraction and modularity**

- Allowing more natural and modular modeling of real-world objects, which are acting concurrently

# Concurrency in Source

- Running concurrent threads using [Source §3 Concurrent](#)
  - Use the primitive function `concurrent_execute`:

`concurrent_execute(  $f_1$ ,  $f_2$ , ...,  $f_n$  )`

- Each  $f_i$  must be a nullary function that returns `undefined`
- The function `concurrent_execute` sets up a separate **thread**  $t_i$  that executes the body of  $f_i$
- These threads  $t_i$  all run **concurrently** with the main thread

# Example 1: Concurrency in Source

```
function sum_1_to_n(n) {  
  let sum = 0;  
  for (let i = 1; i <= n; i = i + 1) { sum = sum + i; }  
  return sum;  
}  
  
function thread1() {  
  display(get_time() - prog_start_tm, "Thread 1 start time:");  
  sum_1_to_n(5000);  
  display(get_time() - prog_start_tm, "Thread 1 middle time:");  
  sum_1_to_n(5000);  
  display(get_time() - prog_start_tm, "Thread 1 end time:");  
}  
  
function thread2() {  
  display(get_time() - prog_start_tm, "Thread 2 start time:");  
  sum_1_to_n(5000);  
  display(get_time() - prog_start_tm, "Thread 2 middle time:");  
  sum_1_to_n(5000);  
  display(get_time() - prog_start_tm, "Thread 2 end time:");  
}  
  
// continue next page
```

[Show in  
Playground](#)



# Example 1: Concurrency in Source

```
// continue from previous page

let prog_start_tm = 0;

// Sequential run
prog_start_tm = get_time();
display("Sequential run:");
thread1(); thread2();

// Concurrent run
prog_start_tm = get_time();
display("Concurrent run:");
concurrent_execute(thread1, thread2);
```

[Show in  
Playground](#)

# Example 1: Concurrency in Source

- Sample output:

## "Sequential run:"

```
Thread 1 start time: 2  
Thread 1 middle time: 88  
Thread 1 end time: 159  
Thread 2 start time: 159  
Thread 2 middle time: 226  
Thread 2 end time: 292
```

## "Concurrent run:"

```
Thread 2 start time: 0  
Thread 1 start time: 0  
Thread 2 middle time: 157  
Thread 1 middle time: 158  
Thread 1 end time: 292  
Thread 2 end time: 294
```

[Show in  
Playground](#)

- What do you observe?
  - “Concurrent run” took about same amount of time as “Sequential run”
    - **Concurrency** does not imply **parallelism**

# Concurrency vs Parallelism

- **Concurrency**

- When execution of multiple tasks **overlap in time**
- Can happen on a single core or processor

- **Parallelism**

- When execution of multiple tasks **occur simultaneously in time**
- Must happen on multiple cores or processors

## Example 2: Interleaving of Threads

- Another example to show the **interleaving** of **events** in concurrent threads

[Show in  
Playground](#)

```
function delay(ms) {
    const start_tm = get_time();
    while (get_time() - start_tm < ms) {
        /* do nothing */
    }
}

function thread_A() {
    delay(20 * math_random());
    display("A1");
    delay(20 * math_random());
    display("A2");
    delay(20 * math_random());
    display("A3");
}
```


```
function thread_B() {
    delay(20 * math_random());
    display("B1");
    delay(20 * math_random());
    display("B2");
    delay(20 * math_random());
    display("B3");
}

// Concurrent run
concurrent_execute(thread_A,
                   thread_B);
```

## Example 2: Interleaving of Threads

- Outputs from a few sample runs:

[Show in  
Playground](#)



## Example 3: Interleaving of Threads

```
function delay(ms) {  
    const start_tm = get_time();  
    while (get_time() - start_tm < ms) {  
        /* do nothing */  
    }  
}
```

- This example shows two concurrent threads updating a **shared variable**

```
function thread_A() {  
    const s = x;  
    const t = s + 1;  
    x = t;  
}  
  
function thread_B() {  
    const s = x;  
    const t = s - 1;  
    x = t;  
}  
  
let x = 0;  
  
concurrent_execute(thread_A,  
                    thread_B);  
  
delay(200);  
display(x, "Final x:");
```

[Show in  
Playground](#)

# Example 3: Interleaving of Threads

- Outputs from a few sample runs:

[Show in  
Playground](#)

Final x: 0

Final x: -1

Final x: 1

- How did we get these different results?
  - We have a **race condition**, where “correct” results are dependent on the ordering or timing of the concurrent events
- Which ones are considered “**correct**”?

# Correct Behavior of Concurrent Programs

- A concurrent program is considered **correct** if it produces the same result **as if** the threads had run **sequentially** in **some** order
  - It does not require the threads *to actually run sequentially*, but only to produce results that are the same **as if they had run sequentially**
  - There may be more than one possible “correct” result produced by the concurrent program, because we require only that the result be the same as for **some** sequential order



# Mechanisms for Controlling Concurrency

- General mechanisms that allow us to **constrain the interleaving** of concurrent threads to **ensure correct** program behavior
- Many mechanisms have been developed
  - Also known as **synchronization mechanisms**
- Here, we consider one of them, the ***serializer***

# Serialization

- **Serialization** implements the following idea:
  - Threads will execute concurrently, but **certain sets of functions** cannot be executed concurrently
- **Serialization** creates **distinguished sets of functions** such that only one execution of a function in each **serialized set** is permitted to happen at a time
  - If some function in the set is being executed, then a thread that attempts to execute any function in the set will be **forced to wait** until the first execution has finished
- We can use **serialization** to control access to **shared variables**

# Serializers in Source

- **Serializers** are constructed by the `make_serializer` function
  - `make_serializer` is not a primitive or pre-declared function in **Source §3 Concurrent**
  - **Source §3 Concurrent** provides additional primitive functions to allow the implementation of `make_serializer`
- A **serializer** takes a function as argument and returns a serialized function that behaves like the original function
- All calls to a given **serializer** return serialized functions in the same set

# Example Use of Serializers

- Executing the following

```
let x = 10;  
const s = make_serializer();  
concurrent_execute(s(() => { x = x * x; }),  
                  s(() => { x = x + 1; }));
```

can produce only two possible values for **x**: 101 or 121

## Example 4: Serializers

```
function make_serializer() {...}

function delay(ms) {...}

const ser = make_serializer();
```

- This example adds the use of a serializer to Example 3
- Its output is now always

**Final x: 0**

```
function thread_A() {
  const s = x;
  const t = s + 1;
  x = t;
}

function thread_B() {
  const s = x;
  const t = s - 1;
  x = t;
}

let x = 0;

concurrent_execute(ser(thread_A),
                   ser(thread_B));

delay(200);
display(x, "Final x:");
```

[Show in  
Playground](#)

# Example 5: Serializers

```
function make_serializer() {...}
function delay(ms) {...}

const ser_x = make_serializer();
const ser_y = make_serializer();

function thread_A() {
  ser_y(() => {
    const s = y;
    const t = s - 1;
    y = t;
  })();

  ser_x(() => {
    const s = x;
    const t = s + 1;
    x = t;
  })();
}
```

```
function thread_B() {
  ser_x(() => {
    const s = x;
    const t = s - 1;
    x = t;
  })();

  ser_y(() => {
    const s = y;
    const t = s + 1;
    y = t;
  })();
}

let x = 0;
let y = 0;
concurrent_execute(thread_A, thread_B);
delay(200);
display(x, "Final x:");
display(y, "Final y:");
```

[Show in  
Playground](#)

Output :

Final x: 0  
Final y: 0

# Implementing Serializers

- We implement serializers in terms of a more primitive synchronization mechanism called a ***mutex***
  - Abbreviation for ***mutual exclusion***
  - A **mutex** is an object that supports two operations:
    - **acquired** and **released**
  - Once a mutex has been acquired, no other acquire operations on that mutex may proceed until the mutex is released
- In our implementation, each **serializer** has an associated **mutex**

# Implementing Serializers

```
function make_serializer() {  
    const mutex = make_mutex();  
    return f => {  
        function serialized_f() {  
            mutex("acquire");  
            const val = f();  
            mutex("release");  
            return val;  
        }  
        return serialized_f;  
    };  
}
```



# Implementing Mutexes

```
function make_mutex() {  
  const cell = list(false);  
  function the_mutex(m) {  
    return m === "acquire"  
      ? test_and_set(cell)  
      ? the_mutex("acquire") // retry  
      : true  
      : m === "release"  
      ? clear(cell)  
      : error(m, "Unknown request -- mutex");  
  }  
  return the_mutex;  
}
```

# Implementing Mutexes

- The `clear` and `test_and_set` functions are primitive functions provided in Source §3 Concurrent
  - Each is performed ***atomically*** (i.e. without interruption)
  - They have the following meanings\*:

```
function clear(cell) {  
    set_head(cell, false);  
}
```

```
function test_and_set(cell) {  
    if (head(cell)) {  
        return true;  
    } else {  
        set_head(cell, true);  
        return false;  
    }  
}
```

# Outline

- Concurrency
- **Applications of Serializers**
- Concurrent programming for multiple shared resources
- More synchronization mechanisms
- Classical synchronization problems

## Example 6: Bank Account

- Want to implement a **bank account** that allows **concurrent withdrawal** and **deposit** transactions
  - For example, a joint bank account concurrently accessed from different ATMs

# Example 6a: Implementation without Serializers

```
function delay(ms) {...}

function make_account(balance) {
  function withdraw(amount) {
    if (balance >= amount) {
      balance = balance - amount;
      return balance;
    } else {
      return "Insufficient funds";
    }
  }
  function deposit(amount) {
    balance = balance + amount;
    return balance;
  }
  function dispatch(m) {
    return m === "withdraw"
      ? withdraw
      : m === "deposit"
      ? deposit
      : m === "balance"
      ? balance
      : error(m,
        "Unknown request");
  }
  return dispatch;
}
```

```
const my_account = make_account(100);

function thread_A() {
  my_account("withdraw")(50);
  my_account("deposit")(100);
}

function thread_B() {
  my_account("withdraw")(50);
  my_account("deposit")(100);
}

concurrent_execute(thread_A, thread_B);

delay(200);
display(my_account("balance"),
  "Final balance:");
```

[Show in  
Playground](#)

# Example 6a: Implementation without Serializers

- Outputs from a few sample runs:

**Final balance: 100**

**Final balance: 150**

**Final balance: 200**

**Final balance: 250**

## Example 6b: Implementation with Serializers

```
function make_serializer() {...}
function delay(ms) {...}

function make_account(balance) {
  function withdraw(amount) {
    if (balance >= amount) {
      balance = balance - amount;
      return balance;
    } else {
      return "Insufficient funds";
    }
  }
  function deposit(amount) {
    balance = balance + amount;
    return balance;
  }

  const protect = make_serializer();
  const protect_withdraw = protect(withdraw);
  const protect_deposit = protect(deposit);
```

```
function dispatch(m) {
  return m === "withdraw"
    ? protect_withdraw
    : m === "deposit"
    ? protect_deposit
    : m === "balance"
    ? balance
    : error(m, "Unknown request");
}
return dispatch;
}

// continue next page
```

[Show in  
Playground](#)

## Example 6b: Implementation with Serializers

```
// continue from previous page

const my_account = make_account(100);

function thread_A() {
  my_account("withdraw")(50);
  my_account("deposit")(100);
}

function thread_B() {
  my_account("withdraw")(50);
  my_account("deposit")(100);
}

concurrent_execute(thread_A, thread_B);

delay(200);
display(my_account("balance"),
        "Final balance:");
```

- Its output is now always

**Final balance: 200**

[Show in  
Playground](#)



# Outline

- Concurrency
- Applications of Serializers
- **Concurrent programming for multiple shared resources**
- More synchronization mechanisms
- Classical synchronization problems

## Example 7: Exchanging Account Balances

- Want to implement a function to **exchange balances** of two bank accounts and can work **concurrently** with **withdrawal**, **deposit** and **other exchange balances** transactions
- A concurrent operation that manipulates **multiple shared resources**

## Example 7a: Implementation (Incorrect)

```
function make_serializer() {...}
function delay(ms) {...}

function make_account(balance) {
  function withdraw(amount) {
    if (balance >= amount) {
      balance = balance - amount;
      return balance;
    } else {
      return "Insufficient funds";
    }
  }
  function deposit(amount) {
    balance = balance + amount;
    return balance;
  }

  const protect = make_serializer();
  const protect_withdraw = protect(withdraw);
  const protect_deposit = protect(deposit);
```

```
function dispatch(m) {
  return m === "withdraw"
    ? protect_withdraw
    : m === "deposit"
    ? protect_deposit
    : m === "balance"
    ? balance
    : error(m,
      "Unknown request");
}
return dispatch;
}

// continue next page
```

[Show in  
Playground](#)

# Example 7a: Implementation (Incorrect)

```
// continue from previous page

function exchange(accounts) {
  const account1 = head(accounts);
  const account2 = tail(accounts);
  const difference = account1("balance")
    - account2("balance");
  account1("withdraw")(difference);
  account2("deposit")(difference);
}

const acc1 = make_account(100);
const acc2 = make_account(200);

function thread_A() {
  exchange(pair(acc1, acc2));
}
function thread_B() {
  exchange(pair(acc1, acc2));
}

concurrent_execute(thread_A, thread_B);

delay(200);
display(acc1("balance"), "Acc.1 final balance:");
display(acc2("balance"), "Acc.2 final balance:");
```

- Outputs from a few sample runs:

Acc.1 final balance: 300  
Acc.2 final balance: 0

Acc.1 final balance: 100  
Acc.2 final balance: 200

Acc.1 final balance: 200  
Acc.2 final balance: 100

[Show in  
Playground](#)

## Example 7b: Implementation (Correct\*)

```
function make_serializer() {...}
function delay(ms) {...}

function make_account_and_serializer(balance) {
  function withdraw(amount) {
    if (balance >= amount) {
      balance = balance - amount;
      return balance;
    } else {
      return "Insufficient funds";
    }
  }
  function deposit(amount) {
    balance = balance + amount;
    return balance;
  }

  const serializer = make_serializer();
```

```
function dispatch(m) {
  return m === "withdraw"
    ? withdraw
    : m === "deposit"
    ? deposit
    : m === "balance"
    ? balance
    : m === "serializer"
    ? serializer
    : error(m,
      "Unknown request");
}
return dispatch;
}

// continue next page
```

[Show in  
Playground](#)

## Example 7b: Implementation (Correct\*)

*// continue from previous page*

```
function exchange(accounts) {  
  const account1 = head(accounts);  
  const account2 = tail(accounts);  
  const difference = account1("balance") - account2("balance");  
  account1("withdraw")(difference);  
  account2("deposit")(difference);  
}
```

```
function serialized_exchange(accounts) {  
  const account1 = head(accounts);  
  const account2 = tail(accounts);  
  const serializer1 = account1("serializer");  
  const serializer2 = account2("serializer");  
  serializer1(serializer2(exchange))(accounts);  
}
```

*// continue next page*

[Show in  
Playground](#)

## Example 7b: Implementation (Correct\*)

```
// continue from previous page

const acc1 = make_account_and_serializer(100);
const acc2 = make_account_and_serializer(200);

function thread_A() {
    serialized_exchange(pair(acc1, acc2));
}
function thread_B() {
    serialized_exchange(pair(acc1, acc2));
}

concurrent_execute(thread_A, thread_B);

delay(200);
display(acc1("balance"), "Acc.1 final balance:");
display(acc2("balance"), "Acc.2 final balance:");
```

- Its output is now always

**Acc.1 final balance: 100**  
**Acc.2 final balance: 200**

[Show in  
Playground](#)

## Example 7c: Beware

- Now let's make the following change to the previous program:

```
// continue from previous page

const acc1 = make_account_and_serializer(100);
const acc2 = make_account_and_serializer(200);

function thread_A() {
    serialized_exchange(pair(acc1, acc2));
}
function thread_B() {
    // serialized_exchange(pair(acc1, acc2));
    serialized_exchange(pair(acc2, acc1)); // changed here
}

concurrent_execute(thread_A, thread_B);

delay(200);
display(acc1("balance"), "Acc.1 final balance:");
display(acc2("balance"), "Acc.2 final balance:");
```

[Show in  
Playground](#)



# Example 7c: Beware

- We get this error message:

[Show in  
Playground](#)

```
Potential infinite loop detected.  
...
```

- Why?
  - Program has run into a **deadlock**
    - **Deadlock** — no thread can make any progress
- How?

## Example 7c: How Deadlock Happened?

- How did the **deadlock** happen?

```
function serialized_exchange(accounts) {  
  const account1 = head(accounts);  
  const account2 = tail(accounts);  
  const serializer1 = account1("serializer");  
  const serializer2 = account2("serializer");  
  serializer1(serializer2(exchange))(accounts);  
}
```

```
function make_serializer() {  
  const mutex = make_mutex();  
  return f => arg => {  
    mutex("acquire");  
    const val = f(arg);  
    mutex("release");  
    return val;  
  };  
}
```

# Example 7c: How Deadlock Happened?

- Original program:

**thread\_A:**

```
acc1_mutex("acquire");  
acc2_mutex("acquire");  
exchange(pair(acc1, acc2));  
acc2_mutex("release");  
acc1_mutex("release");
```

**thread\_B:**

```
acc1_mutex("acquire");  
acc2_mutex("acquire");  
exchange(pair(acc1, acc2));  
acc2_mutex("release");  
acc1_mutex("release");
```

- Modified program:

**thread\_A:**

```
acc1_mutex("acquire");  
acc2_mutex("acquire");  
exchange(pair(acc1, acc2));  
acc2_mutex("release");  
acc1_mutex("release");
```

**thread\_B:**

```
acc2_mutex("acquire");  
acc1_mutex("acquire");  
exchange(pair(acc2, acc1));  
acc1_mutex("release");  
acc2_mutex("release");
```

# Example 7c: How Deadlock Happened?

- Modified program:

thread\_A:

```
acc1_mutex("acquire");  
acc2_mutex("acquire");  
exchange(pair(acc1, acc2));  
acc2_mutex("release");  
acc1_mutex("release");
```

thread\_B:

```
acc2_mutex("acquire");  
acc1_mutex("acquire");  
exchange(pair(acc2, acc1));  
acc1_mutex("release");  
acc2_mutex("release");
```

- Possible interleaving that led to deadlock:

Time	thread_A	thread_B
1	acc1_mutex("acquire");	
2		acc2_mutex("acquire");
3	acc2_mutex("acquire");	
4	<i>blocked</i>	acc1_mutex("acquire");
5	<i>blocked</i>	<i>blocked</i>
⋮	⋮	⋮

# Deadlock Avoidance

- How to avoid deadlock (for this situation)?
  - Give each account a **unique identification**
  - Each concurrent thread will **protect the lowest-numbered account first**
- There are other situations that require more **sophisticated deadlock-avoidance techniques**, or where **deadlock cannot be avoided** at all

# Outline

- Concurrency
- Applications of Serializers
- Concurrent programming for multiple shared resources
- **More synchronization mechanisms**
- Classical synchronization problems

# Synchronization Mechanisms

- Synchronization mechanisms
  - **Serializers**
  - **Mutexes**
    - A **mutex** can be used to protect a **critical section** of a thread that accesses a **shared resource**
    - It guarantees **mutually exclusive** access to the **shared resource** at any given time
    - Threads that want to access a resource protected by a mutex must **wait until** the currently active thread is finished and **unlock** the mutex
  - **Semaphores**
  - etc.

# Semaphores

- **Semaphore**

- A generalized synchronization mechanism
- Only behaviors are specified; can have different implementations
- Provides
  - A way to **block** a number of processes/threads
    - Known as **sleeping process/threads**
  - A way to **unblock/wake up** one or more sleeping processes/threads
- Proposed by **Edgar W. Dijkstra** in 1965



# Semaphores

- A semaphore **S** contains an **integer** value
  - Initialized to any **non-negative** values initially
- Two **atomic** semaphore operations:

**Down:** (a.k.a. **P** and **Wait**)

1. If **S**  $\leq 0$ , blocks (go to sleep)
2. Decrement **S**

**Up:** (a.k.a. **V** and **Signal**)

1. Increment **S**
2. Wakes up one sleeping process if any

# Semaphores: Implementation using Mutex\*

```
function make_semaphore(val) {  
  const val_mutex = make_mutex();  
  function semaphore(m) {  
    if (m === "up") {  
      val_mutex("acquire");  
      val = val + 1;  
      // Wake up one waiting thread if any.  
      val_mutex("release");  
    } else if (m === "down") {  
      val_mutex("acquire");  
      while (val <= 0) {  
        val_mutex("release");  
        // Put thread to sleep.  
        val_mutex("acquire");  
      }  
      val = val - 1;  
      val_mutex("release");  
    } else {  
      error(m, "Unknown request -- semaphore");  
    }  
  }  
  return semaphore;  
}
```

[Show in  
Playground](#)

# Outline

- Concurrency
- Applications of Serializers
- Concurrent programming for multiple shared resources
- More synchronization mechanisms
- **Classical synchronization problems**

# Classical Synchronization Problems

- The Producers-Consumers problem
- The Readers-Writers problem
- The Dining Philosophers problem

# The Producers-Consumers Problem

- The **Producers-Consumers** problem specification
  - Threads share a **bounded circular buffer** (array) of size **N**
  - **Producers** produce items and insert them to buffer
    - Only when buffer is **not full** (  $< N$  items)
  - **Consumers** remove items from buffer
    - Only when buffer is **not empty** (  $> 0$  items)

# The Producers-Consumers Problem: A Solution

```
const N = ...;  
const buffer = [];  
let count = 0;  
let in_pos = 0;  
let out_pos = 0;
```

```
const mutex = make_semaphore(1);  
const not_full = make_semaphore(N);  
const not_empty = make_semaphore(0);
```

```
function producer() {  
    while (true) {  
        const item = produce_item(...);  
  
        not_full("down");  
        mutex("down");  
  
        buffer[in_pos] = item;  
        in_pos = (in_pos + 1) % N;  
        count = count + 1;  
  
        mutex("up");  
        not_empty("up");  
    }  
}
```

```
function consumer() {  
    while (true) {  
        not_empty("down");  
        mutex("down");  
  
        const item = buffer[out_pos];  
        out_pos = (out_pos + 1) % N;  
        count = count - 1;  
  
        mutex("up");  
        not_full("up");  
  
        consume_item(item);  
    }  
}
```

# The Readers-Writers Problem

- The **Readers-Writers** problem specification
  - Threads share a data structure **D**
  - **Readers** retrieve information from **D**
    - Each reader can access **D** with other readers
  - **Writers** modify information in **D**
    - Each writer must have exclusive access to **D**

# The Readers-Writers Problem: A Solution\*

```
const D = make_data_structure(...);

let num_readers = 0;

const mutex = make_semaphore(1);
const room_empty = make_semaphore(1);

function writer() {
  while (true) {
    room_empty("down");

    modify_data(D, ...);

    room_empty("up");
  }
}
```

```
function reader() {
  while (true) {
    mutex("down");

    num_readers = num_readers + 1;
    if (num_readers === 1) {
      room_empty("down");
    } else {}

    mutex("up");

    read_data(D, ...);

    mutex("down");


    num_readers = num_readers - 1;
    if (num_readers === 0) {
      room_empty("up");
    } else {}

    mutex("up");
  }
}
```



# The Dining Philosophers Problem

- The **Dining Philosophers** problem specification
  - Five **philosophers** are seated at a round table
    - With a big bowl of noodles at the center
  - Five single **chopsticks** are placed between each pair of adjacent philosophers
  - Each philosopher must alternately **think** and **eat**
  - A philosopher can only eat when he/she has **both left and right chopsticks**
  - Assume unlimited supply of noodles and unlimited appetite



[Image from [www.geeksforgeeks.org](http://www.geeksforgeeks.org)]

# The Dining Philosophers Problem: Attempt #1

```
function philosopher(i) {  
    while (true) {  
        think();  
        take_left_chopstick(i);  
        take_right_chopstick(i);  
        eat();  
        put_left_chopstick(i);  
        put_right_chopstick(i);  
    }  
}
```

- Any problem?
  - **Deadlock**
    - E.g. all philosophers simultaneously take up left chopsticks, and none can proceed

# The Dining Philosophers Problem: Attempt #2

```
function philosopher(i) {  
    while (true) {  
        think();  
        mutex("acquire");  
        take_left_chopstick(i);  
        take_right_chopstick(i);  
        eat();  
        put_left_chopstick(i);  
        put_right_chopstick(i);  
        mutex("release");  
    }  
}
```

- Any problem?
  - Only **one philosopher** can eat at a time
  - Possible **starvation** of some philosophers

# The Dining Philosophers Problem: A Good Solution

- An exercise or research for you

# Summary

- Applications of Serializers
- Concurrent programming for multiple shared resources
  - Be careful of deadlocks
- More synchronization mechanisms
  - Mutexes
  - Semaphores
- Classical synchronization problems
  - The Producers-Consumers problem
  - The Readers-Writers problem
  - The Dining Philosophers problem