

Discussion Group Problems for Week 9

For: March 14–March 18

**Problem 1. Quadratic Probing**

Quadratic probing is another open-addressing scheme very similar to linear probing. Recall that a linear probing implementation searches the next bucket on a collision.

We can also express linear probing with the following pseudocode (on insertion of element  $x$ ):

```
for i in 0..m:
    if buckets[hash(x) + i % m] is empty:
        insert x into this bucket
        break
```

Quadratic probing follows a very similar idea. We can express it as follows:

```
for i in 0..m:
    // increment by squares instead
    if buckets[hash(x) + i * i % m] is empty:
        insert x into this bucket
        break
```

- (a) Consider a hash table with size 7 with hash function  $h(x) = x \% 7$ . We insert the following elements in the order given: 5, 12, 19, 26, 2. What does the final hash table look like?
- (b) Continuing from the above question, we now delete the following elements in the order given: 12, 5. What does the final hash table look like?
- (c) Can you construct a case where quadratic probing fails to insert an element despite the table not being full?

**Problem 2. Table Resizing**

Suppose we follow these rules for an implementation of an open-addressing hash table, where  $n$  is the number of items in the hash table and  $m$  is the size of the hash table.

- (a) If  $n = m$ , then the table is *quadrupled* (resize  $m$  to  $4m$ )
- (b) If  $n < m/4$ , then the table is *shrunk* (resize  $m$  to  $m/2$ )

What is the minimum number of insertions between 2 resize events? What about deletions?

**Problem 3. Implementing Union/Intersection of Sets**

Consider the following implementations of sets. How would intersect and union be implemented for each of them?

- (a) Hash table with open addressing
- (b) Hash table with chaining

#### Problem 4. Binary Counter

Binary counter ADT is a data structure that counts in base two, i.e. 0s and 1s. Binary Counter ADT supports two operations:

- `increment()` increases the counter by 1
- `read()` reads the current value of the binary counter

To make it clearer, suppose that we have a  $k$ -bit binary counter. Each bit is stored in an array  $A$  of size  $k$ , where  $A[k]$  denotes the  $k$ -th bit (0-th bit denotes the least significant bit). For example, suppose  $A = [1, 1, 0]$ , which corresponds to the number 011 in binary. Calling `increment()` will yield  $A = [0, 0, 1]$ , i.e. 100. Calling `increment()` again will yield  $A = [1, 0, 1]$ , the number 101 in binary.

Suppose that the  $k$ -bit binary counter starts at 0, i.e. all the values in  $A$  is 0. A loose bound on the time complexity if `increment()` is called  $n$  times is  $O(nk)$ . What is the amortized time complexity of `increment()` operation if we call `increment()`  $n$  times?

#### Problem 5. Stack 2 Queue

Do you know that we actually can implement a queue using two stacks? But is it really efficient?

- (a) Design an algorithm to push and pop an element from the queue.
- (b) Determine the *worst case* and *amortized* runtime for each operation.

#### Problem 6. Scapegoat Trees

Consider the Scapegoat Tree data structure that we have implemented in Problem Sets 4-5. We assume that only insertions are performed on our Scapegoat Tree. In this question, we will use amortized analysis to reason about the performance of inserts on Scapegoat Trees.

- (a) Suppose we are about to perform the `rebuild` operation on a node  $v$ . Show that the amount of entries that *must* have been inserted into node  $v$  since it was **last rebuilt** is  $\Omega(\text{size}(v))$ .
- (b) Show that the *depth* of any insertion is  $O(\log n)$
- (c) Now use the previous two parts to show that the amortized cost of an insertion in a Scapegoat Tree is  $O(\log n)$ . *Hint:* Suppose a constant amount is “deposited” at every node traversed on an insertion.

**Problem 7. Tabulation Hashing**

Suppose we are creating a hash function keys  $N = \log n$  bits long, mapped to buckets indexed by  $M = \log m$ . So, the hash function maps an  $N$  bit key to an  $M$  bit identifier for a bucket. To do this, we construct a 2D-array  $T[2, N]$  and fill each entry in the table with a random  $M$  bit value.

Now to hash a key, we just XOR the key and the table:

```
hash = 0
for (j = 1 to N)
    hash = hash XOR T[key[j], j]
```

Note that `key[j]` denotes the  $j$ -th bit of the key.

**Problem 7.a.** Show that for a given key  $k$  and bucket  $b$ ,  $\Pr[h(k) = b] = 1/m$ , and that for two keys  $k_1 \neq k_2$ ,  $\Pr[h(k_1) = h(k_2)] \leq 1/m$

**Problem 7.b.** How much space does this table require?

Now let's generalize the function. Choose some integer  $R$  that divides  $N$ . (In the previous example,  $R = 1$ .) Now generate a 2D table of size  $[2^R, N/R]$ . As before, fill each entry with a random  $M$  bit value. As before, take the hash by breaking the key into  $R$  bit chunks, look up each chunk in the table, and perform XOR:

```
hash = 0
for (j = 1 to N/R)
    hash = hash XOR T[key[(j-1)R .. jR]][j]
```

Notice now the table needs one entry for each of the  $2^R$  possible chunks of the key.

**Problem 7.c.** What is the size of the table?

**Problem 7.d.** *Bonus* Here's another simple hashing scheme:

Construct a 2D binary array  $A[M, N]$ , where each entry is a random 0 or 1. For key  $k$ , let  $h = Ak$  (think of it as matrix multiplication, where  $k$  is a column vector,  $A$  is a matrix, and the result is an  $m$ -bit column vector).

Is this the same as tabulation hashing?